


CSC242: Introduction to Artificial Intelligence

Lecture 1.3

Please put away all electronic devices


Search Strategies

Uninformed



No additional information
about states

Informed
(Heuristic)



Can identify “promising”
states

Search Strategies

Uninformed

	BFS	DFS (tree)	IDS	Greedy	A*	IDA*
Complete ?	✓	✗	✓	✗	✓ [†]	✓ [†]
Optimal ?	✓ [*]	✗	✓ [*]	✗	✓ [†]	✓ ^{*†}
Time	$O(b^d)$	$O(b^m)$	$O(b^d)$	$O(b^m)$	$O(b^{\epsilon d})$	$O(b^{\epsilon d})$
Space	$O(b^d)$	$O(bm)$	$O(bd)$	$O(b^m)$	$O(b^d)$	$O(bd)$

* If step costs are identical

† With an admissible heuristic

Search Strategies

Uninformed

Informed

	BFS	DFS (tree)	IDS	Greedy	A*	IDA*
Complete ?	✓	✗	✓	✗	✓ [†]	✓ [†]
Optimal ?	✓ [*]	✗	✓ [*]	✗	✓ [†]	✓ ^{*†}
Time	$O(b^d)$	$O(b^m)$	$O(b^d)$	$O(b^m)$	$O(b^{\epsilon d})$	$O(b^{\epsilon d})$
Space	$O(b^d)$	$O(bm)$	$O(bd)$	$O(b^m)$	$O(b^d)$	$O(bd)$


* If step costs are identical

† With an admissible heuristic


Evaluation function

$$f(n) = g(n) + h(n)$$

True cost of path from
start node to node n



Estimated cost of cheapest
path from n to a goal node



Admissible Heuristic

Never **over**estimates the true cost
of a solution

$$f(n) = g(n) + h_{SLD}(n)$$

Search Strategies

Uninformed

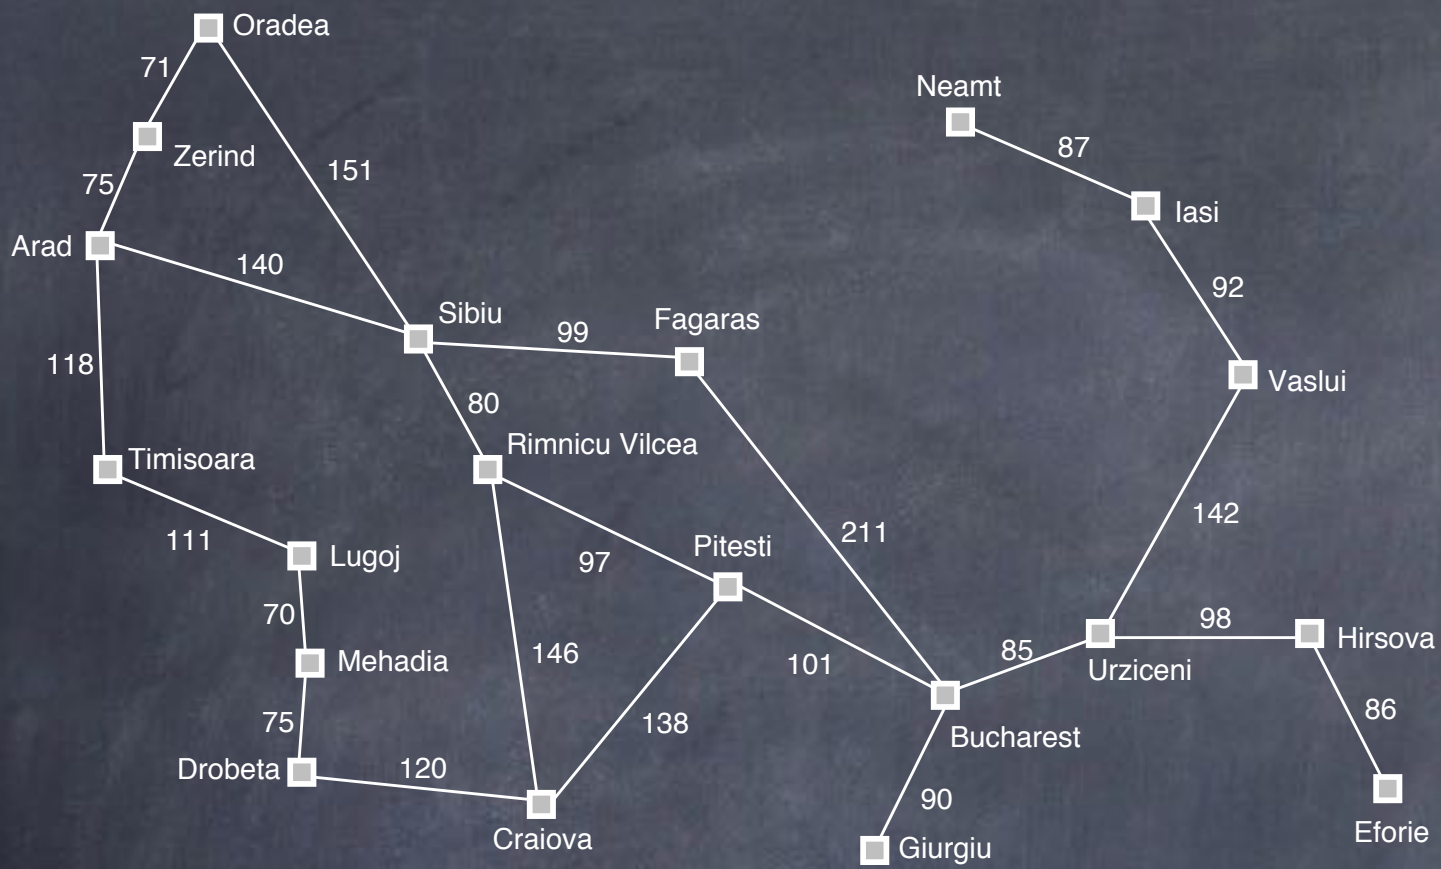
Informed


	BFS	DFS (tree)	IDS	Greedy	A*	IDA*
Complete ?	✓	✗	✓	✗	✓ [†]	✓ [†]
Optimal ?	✓ [*]	✗	✓ [*]	✗	✓ [†]	✓ ^{*†}
Time	$O(b^d)$	$O(b^m)$	$O(b^d)$	$O(b^m)$	$O(b^{\epsilon d})$	$O(b^{\epsilon d})$
Space	$O(b^d)$	$O(bm)$	$O(bd)$	$O(b^m)$	$O(b^d)$	$O(bd)$

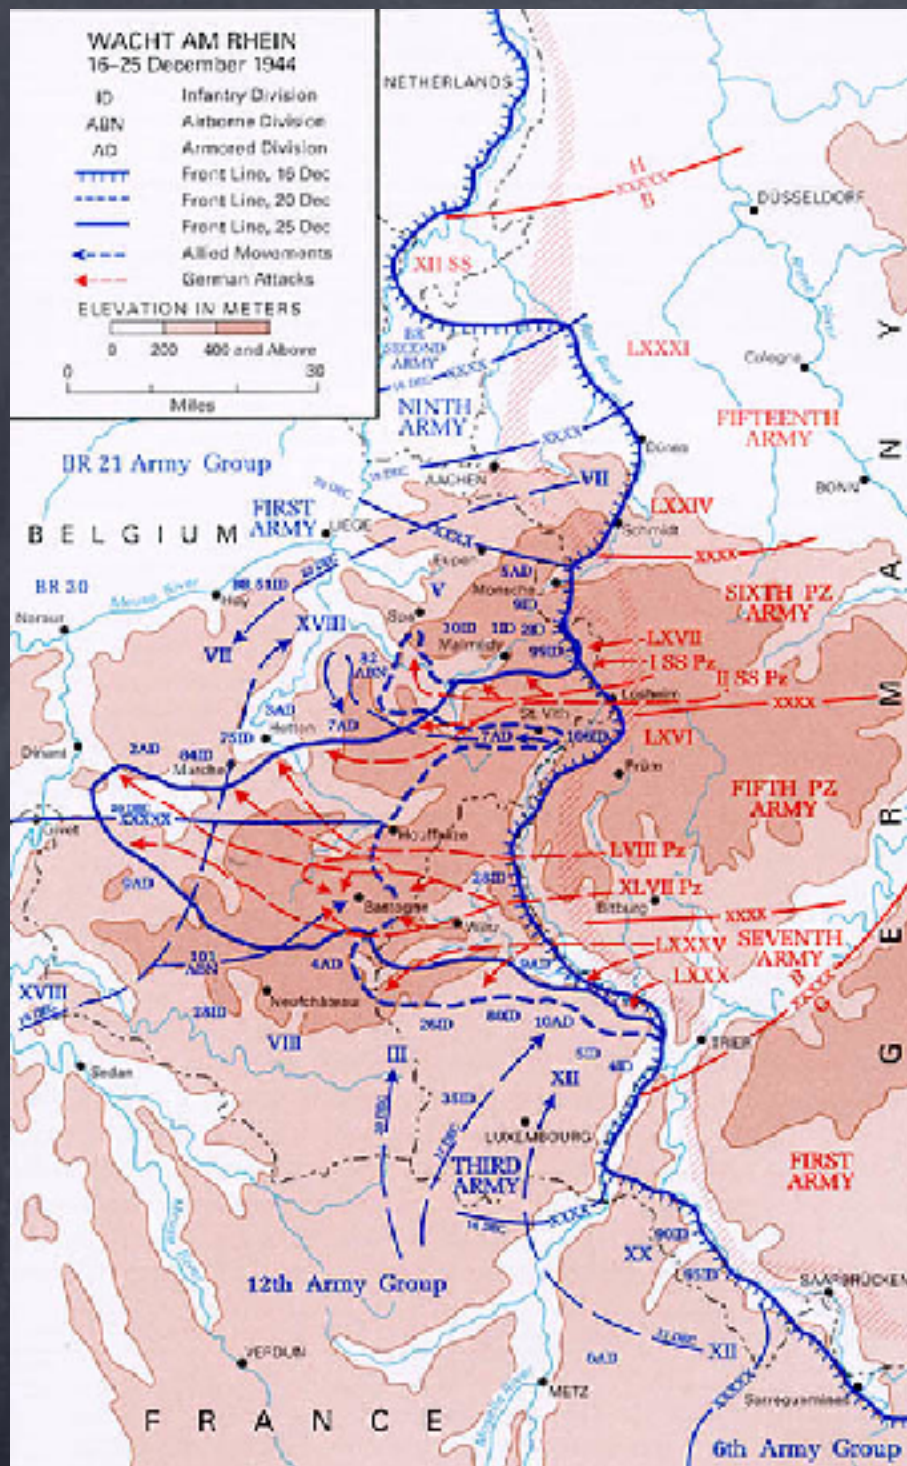
* If step costs are identical

† With an admissible heuristic

Adversarial Search







Battle of The Bulge
16-25 Dec 1944



Multi-Agent Environments

- Unpredictability of other agents => contingencies (strategies)
- Agents goals in conflict => competition

Games

Games



Games

- “Games require the ability to make some decision even when calculating the optimal decision is infeasible”
- “[Games] penalize inefficiency severely”

Games \neq Toy Problems

Games \neq Toy Problems



8-puzzle: 181,440

15-puzzle: $\sim 1.3 \times 10^{12}$

24-puzzle: $\sim 10^{25}$

Games \neq Toy Problems



8-puzzle: 181,440

$$35^{100} = 10^{154}$$

15-puzzle: $\sim 1.3 \times 10^{12}$ (only 10^{40} distinct)

24-puzzle: $\sim 10^{25}$

Games \neq Toy Problems



8-puzzle: 181,440

$$35^{100} = 10^{154}$$

$$2 \times 10^{170}$$

15-puzzle: $\sim 1.3 \times 10^{12}$ (only 10^{40} distinct)

24-puzzle: $\sim 10^{25}$

Abstract Games

Abstract Games

- Deterministic (no chance)
- Nondeterministic (dice, cards, etc.)

Abstract Games

- Perfect information (fully observable)
- Imperfect information (partially observable)

Abstract Games




- Zero-sum (total payoff the same in any game)
 - What's good for me is bad for you and vice-versa
- Arbitrary utility function

Types of Games


Deterministic (no chance)	Nondeterministic (dice, cards, etc.)
Perfect information (fully observable)	Imperfect information (partially observable)
Zero-sum (total payoff the same in any game)	Arbitrary utility functions

Outline




- In deterministic, perfect information, zero-sum games:
 - How to find optimal moves
 - How to find good moves when time is limited
- Next time: nondeterministic
- Later: imperfect information, utilities



Terminal States



Next Player: O




X

O

X

O



Term.




X

O

X

O



$O(n!)$

Term.




X

O

X

O




Term.




Deciding what to do in a game requires thinking about what the opponent will do, and having a **strategy** that takes that into account



Not simply a sequence of actions, but a contingency plan that specifies what to do depending on what state one finds oneself in (AIMA 4.3)




Term.	<div><div>X</div><div>O</div><div>X</div></div> <div><div></div><div>O</div><div>X</div></div> <div><div></div><div></div><div>O</div></div>	<div><div>X</div><div>O</div><div>X</div></div> <div><div>O</div><div>O</div><div>X</div></div> <div><div>X</div><div>X</div><div>O</div></div>	<div><div>X</div><div>O</div><div>X</div></div> <div><div></div><div>X</div><div></div></div> <div><div>X</div><div>O</div><div>O</div></div>
Utility	-1	0	+1




Term.	<div> <div>X</div> <div>O</div> <div>X</div> <div></div> <div>O</div> <div>X</div> <div></div> <div>O</div> <div></div> </div>	<div> <div>X</div> <div>O</div> <div>X</div> <div>O</div> <div>O</div> <div>X</div> <div>X</div> <div>X</div> <div>O</div> </div>	<div> <div>X</div> <div>O</div> <div>X</div> <div></div> <div>X</div> <div></div> <div>X</div> <div>O</div> <div>O</div> </div>
Utility	-1	0	+1

MAX

Term.
Utility






MAX

MIN

MAX

Term.

Utility



Minimax Algorithm

$\text{MINIMAX}(s) =$


$\text{UTILITY}(s)$ if $\text{TERMINAL-TEST}(s)$

$\max_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a))$ if $\text{PLAYER}(s) = \text{MAX}$

$\min_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a))$ if $\text{PLAYER}(s) = \text{MIN}$


MAX

MIN




MAX

MIN




MAX

MIN




MAX

MIN




MAX

MIN



MAX

MIN



Minimax Analysis



Time Complexity

$$O(b^m)$$

Space Complexity

“for real games, the time cost is
totally impractical”

Minimax Summary

- Computes the optimal move assuming opponent also plays optimally (i.e., worst-case outcome)
- Explores game tree depth-first all the way to terminal states (end of game)
- Backs up utility values through alternating MIN and MAX (what's best for me is worst for you, and vice-versa)

Minimax Code

AIMA 5.2.1 and Figure 5.3 (page 166)

Not bad, but...


MAX

MIN

MAX

Term.

Utility



A Problem for Minimax

- You can't search all the way to the terminal nodes
- You can't evaluate a node (state) unless you're at a terminal node
- Utility function is defined on terminal states


Imperfect Real-Time Decisions

MAX

MIN

MAX

Term.




MAX

MIN

MAX

Term.



Minimax Algorithm

$\text{MINIMAX}(s) =$

$\text{UTILITY}(s)$ if $\text{TERMINAL-TEST}(s)$

$\max_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a))$ if $\text{PLAYER}(s) = \text{MAX}$

$\min_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a))$ if $\text{PLAYER}(s) = \text{MIN}$

Heuristic Minimax

H-MINIMAX(s) =

$h(s)$

if CUTOFF-TEST(s)

$\max_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a))$ if $\text{PLAYER}(s) = \text{MAX}$

$\min_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a))$ if $\text{PLAYER}(s) = \text{MIN}$

H-Minimax: Cutoff

- When to cutoff search?
 - Time
 - Depth
 - “Quiescence”
 - Can adjust dynamically
 - AIMA 5.4.2
- Combine with iterative deepening

H-Minimax: Heuristic

- How to evaluate non-terminal state?
 - AIMA 5.4.1
 - Chess example
 - Material value
 - Weighted sum
 - Strategic considerations
- Classic time-quality tradeoff

Heuristic Minimax

- Cutoff search before reaching terminal nodes (time, depth, "quiescence")
- Use heuristic evaluation function to estimate state utility
- Backs up utility values through alternating MIN and MAX (what's best for me is worst for you, and vice-versa)

Adversarial Search:

MINIMAX

MINIMAX

- Searches to terminal nodes
- Uses utility function

H-MINIMAX

- Cuts off search before terminals
- Uses heuristic function

Backs up utility values through alternating MIN and MAX (zero-sum game)

For next time:

Chapter 5.3–5.4.2;
5.5–5.6; 5.7–5.9 fyi