

# CSC242: Introduction to Artificial Intelligence

## Lecture 1.4

Please put away all electronic devices

# Adversarial Search



# Games $\neq$ Toy Problems



8-puzzle: 181,440

$$35^{100} = 10^{154}$$

$$2 \times 10^{170}$$

15-puzzle:  $\sim 1.3 \times 10^{12}$  (only  $10^{40}$  distinct)

24-puzzle:  $\sim 10^{25}$

# Types of Games

Deterministic (no chance)	Nondeterministic (dice, cards, etc.)
Perfect information (fully observable)	Imperfect information (partially observable)
Zero-sum (total payoff the same in any game)	Arbitrary utility functions



Deciding what to do in a game requires thinking about what the opponent will do, and having a **strategy** that takes that into account



Not simply a sequence of actions, but a contingency plan that specifies what to do depending on what state one finds oneself in (AIMA 4.3)

# Minimax Algorithm

$\text{MINIMAX}(s) =$

$\text{UTILITY}(s)$  if  $\text{TERMINAL-TEST}(s)$

$\max_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a))$  if  $\text{PLAYER}(s) = \text{MAX}$

$\min_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a))$  if  $\text{PLAYER}(s) = \text{MIN}$



# Minimax Algorithm

$\text{MINIMAX}(s) =$

$\text{UTILITY}(s)$


if  $\text{TERMINAL-TEST}(s)$

$\max_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a))$  if  $\text{PLAYER}(s) = \text{MAX}$

$\min_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a))$  if  $\text{PLAYER}(s) = \text{MIN}$

MAX

MIN





# Minimax Summary

- Computes the optimal move assuming opponent also plays optimally (i.e., worst-case outcome)
- Explores game tree depth-first all the way to terminal states (end of game)
- Backs up utility values through alternating MIN and MAX (what's best for me is worst for you, and vice-versa)

# Minimax Analysis



Time Complexity

$$O(b^m)$$

Space Complexity

“for real games, the time cost is  
totally impractical”



# Heuristic Minimax

$$\text{H-MINIMAX}(s) =$$

$$h(s) \quad \text{if CUTOFF-TEST}(s)$$

$$\max_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) \quad \text{if PLAYER}(s) = \text{MAX}$$

$$\min_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) \quad \text{if PLAYER}(s) = \text{MIN}$$

# H-Minimax

- When to cutoff search?
  - Time, depth, “quiescence”, ...
- How evaluate non-terminal states?
  - Depends on game
  - Tradeoff between time and quality



# Adversarial Search

- In deterministic, perfect information, zero-sum games:
  - How to find optimal moves
    - MINIMAX (requires utility fn)
  - How to find good moves when time is limited
    - H-MINIMAX (requires cutoff &  $h$  fn)





# Google/DeepMind AlphaGo

- Monte Carlo tree search (MCTS) with:
  - Move selection ("Policy")
  - Position evaluation ("Value")
- Neural networks compute mapping from board positions to moves and evaluations
- Trained using combination of supervised and reinforcement learning

# Google/DeepMind AlphaGo

- Monte Carlo tree search (MCTS) with:
  - Move selection ("Policy")
  - Position evaluation ("Value")
- Neural networks compute mapping from board positions to moves and evaluations
- Trained using combination of supervised and reinforcement learning



# Google/DeepMind AlphaGo

- Monte Carlo tree search (MCTS) with:
  - Move selection ("Policy")
  - Position evaluation ("Value")
- Neural networks compute mapping from board positions to moves and evaluations
- Trained using combination of supervised and reinforcement learning

# Google/DeepMind AlphaGo

- Monte Carlo tree search (MCTS) with:
  - Move selection ("Policy") → reduces  $b$
  - Position evaluation ("Value")
- Neural networks compute mapping from board positions to moves and evaluations
- Trained using combination of supervised and reinforcement learning



# Google/DeepMind AlphaGo

- Monte Carlo tree search (MCTS) with:
  - Move selection ("Policy")
  - Position evaluation ("Value")  $\rightarrow h(n)$
- Neural networks compute mapping from board positions to moves and evaluations
- Trained using combination of supervised and reinforcement learning

# Google/DeepMind AlphaGo

- Monte Carlo tree search (MCTS) with:
  - Move selection ("Policy")
  - Position evaluation ("Value")
- Neural networks compute mapping from board positions to moves and evaluations
- Trained using combination of supervised and reinforcement learning



# Google/DeepMind AlphaGo

<http://deepmind.com/alpha-go.html>

Silver, D., et al. (2016). Mastering the game of Go with deep neural networks and tree search. Nature 529, 484–489.


<http://doi.org/10.1038/nature16961>

# Minimax Analysis



$$O(b^m)$$

Time Complexity




$$O(bm)$$

Space Complexity




MAX

MIN



MAX

MIN



$\alpha, \beta$

Highest utility seen so far  
(lower bound on max value)

Lowest utility seen so far  
(upper bound on min value)



MAX

$-\infty, +\infty$

A

$-\infty, +\infty$

MIN

B



3

MAX

$-\infty, +\infty$

A

~~$-\infty, +\infty$~~

B

MIN



3



MAX

$-\infty, +\infty$

A

$-\infty, 3$

MIN

B



3

12

8

MAX

$-\infty, +\infty$

A

~~$-\infty, 33$~~

MIN

B



3

12

8



MAX

~~-39, +100~~

A

3, 3

MIN

B




3

12


8

MAX

MIN








MAX


MIN






MAX

MIN




MAX

MIN





# Alpha-Beta Pruning



AIMA Fig 5.6

# Alpha-Beta Pruning

- During MINIMAX search
  - Keep track of
    - $\alpha$ : value of best choice so far for MAX (lower bound on MAX utility)
    - $\beta$ : value of best choice so far for MIN (upper bound on MIN utility)
  - Prune when value of node is known to be worse than  $\alpha$  (for MAX) or  $\beta$  (for MIN)



# Alpha-Beta Pruning

- Still MINIMAX search
  - Optimal (if you search to the terminals)
  - Optimal with respect to your heuristic function otherwise

# Alpha-Beta Pruning

- Still MINIMAX search
  - Optimal (if you search to the terminals)
  - Optimal with respect to your heuristic function otherwise

Hang on...



# Alpha-Beta Pruning

- Still MINIMAX search
  - Optimal (if you search to the terminals)
  - Optimal with respect to your heuristic function otherwise


DFS with utility function

# Alpha-Beta Pruning Analysis




MAX

MIN



MAX

MIN





# Alpha-Beta Pruning Analysis

Ideal case:

Explore the best successor first:  $O(b^{m/2})$

Branching factor:  $b^{1/2} = \sqrt{b}$

Explore twice as deep a tree in same time

# Alpha-Beta Pruning Analysis

Random case:

Explore successors in random order:  $O(b^{\frac{3}{4}m})$

Branching factor:  $b^{3/4}$

Explore 1/3 deeper tree in same time



# Alpha-Beta Pruning Analysis

“Smart” case:

- e.g., for chess order by:
  - Captures
  - Threats of captures
  - Forward moves
  - Backward moves

# Alpha-Beta Summary

- Easy bookkeeping modification of basic MINIMAX algorithm
- Not hard to come up with “useful” node orderings
- Even random gets you 33% deeper search
- Works with other ways of improving game tree search



# Types of Games

Deterministic (no chance)	Nondeterministic (dice, cards, etc.)
Perfect information (fully observable)	Imperfect information (partially observable)
Zero-sum (total payoff the same in any game)	Arbitrary utility functions


# Stochastic and Partially Observable Games



# Non-deterministic (Stochastic) Games

- A player's possible moves depend on chance (random) elements, e.g., dice

# Stochastic Games





# Stochastic Games

- A player's possible moves depend on chance (random) elements, e.g., dice
- Can't build game tree since don't know what future legal moves will be

MAX


CHANCE

MIN

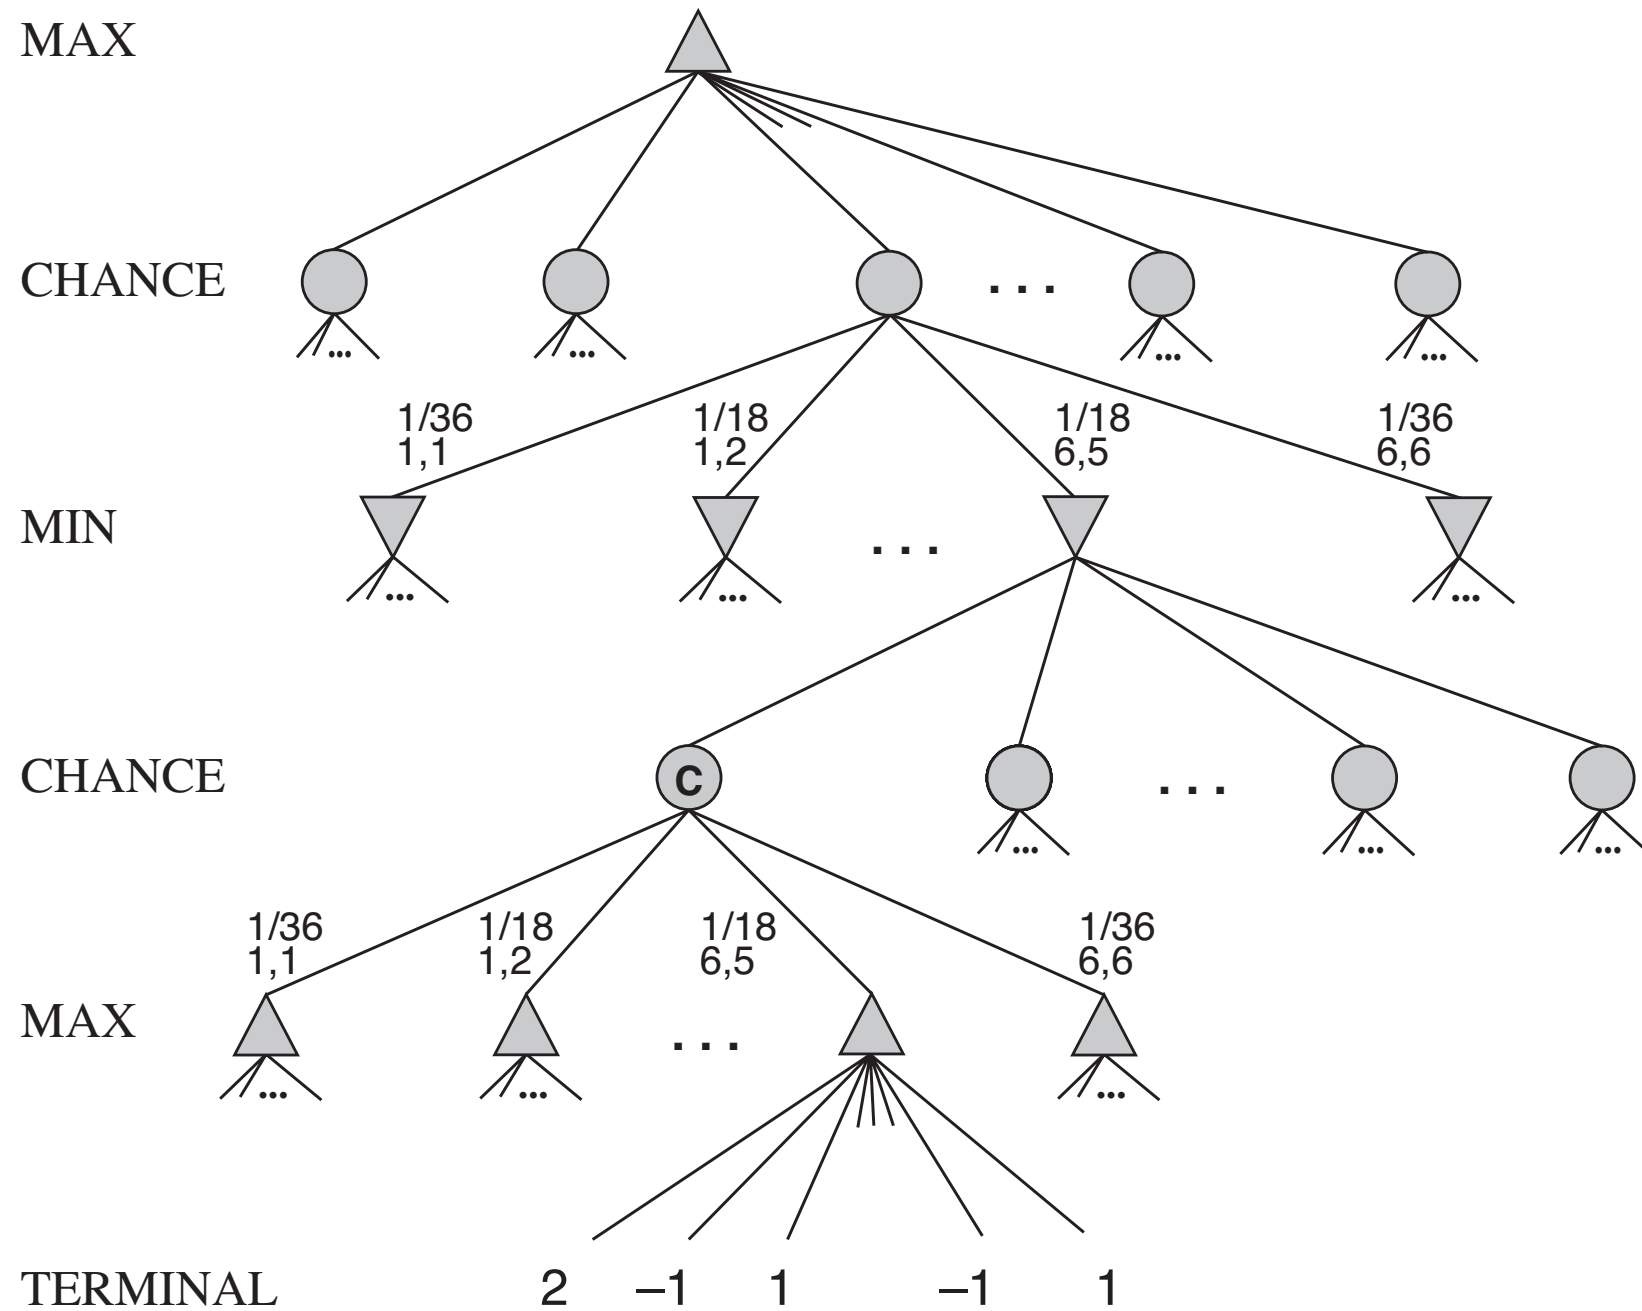
CHANCE

MAX


TERMINAL









MINIMAX?









Average over possible outcomes











# Expectation (Expected Value)

- Weighted average of possibilities
- Sum of the possible outcomes weighted by the likelihood of their occurrence
- What you would expect to win in the long run

# Expecti-Minimax

- Same as MINIMAX for MIN and MAX nodes
- Same backing up utilities from terminal nodes
- Take expectation over chance nodes
  - Weighted average of possible outcomes



MAX


CHANCE

MIN

CHANCE

MAX

TERMINAL



MAX

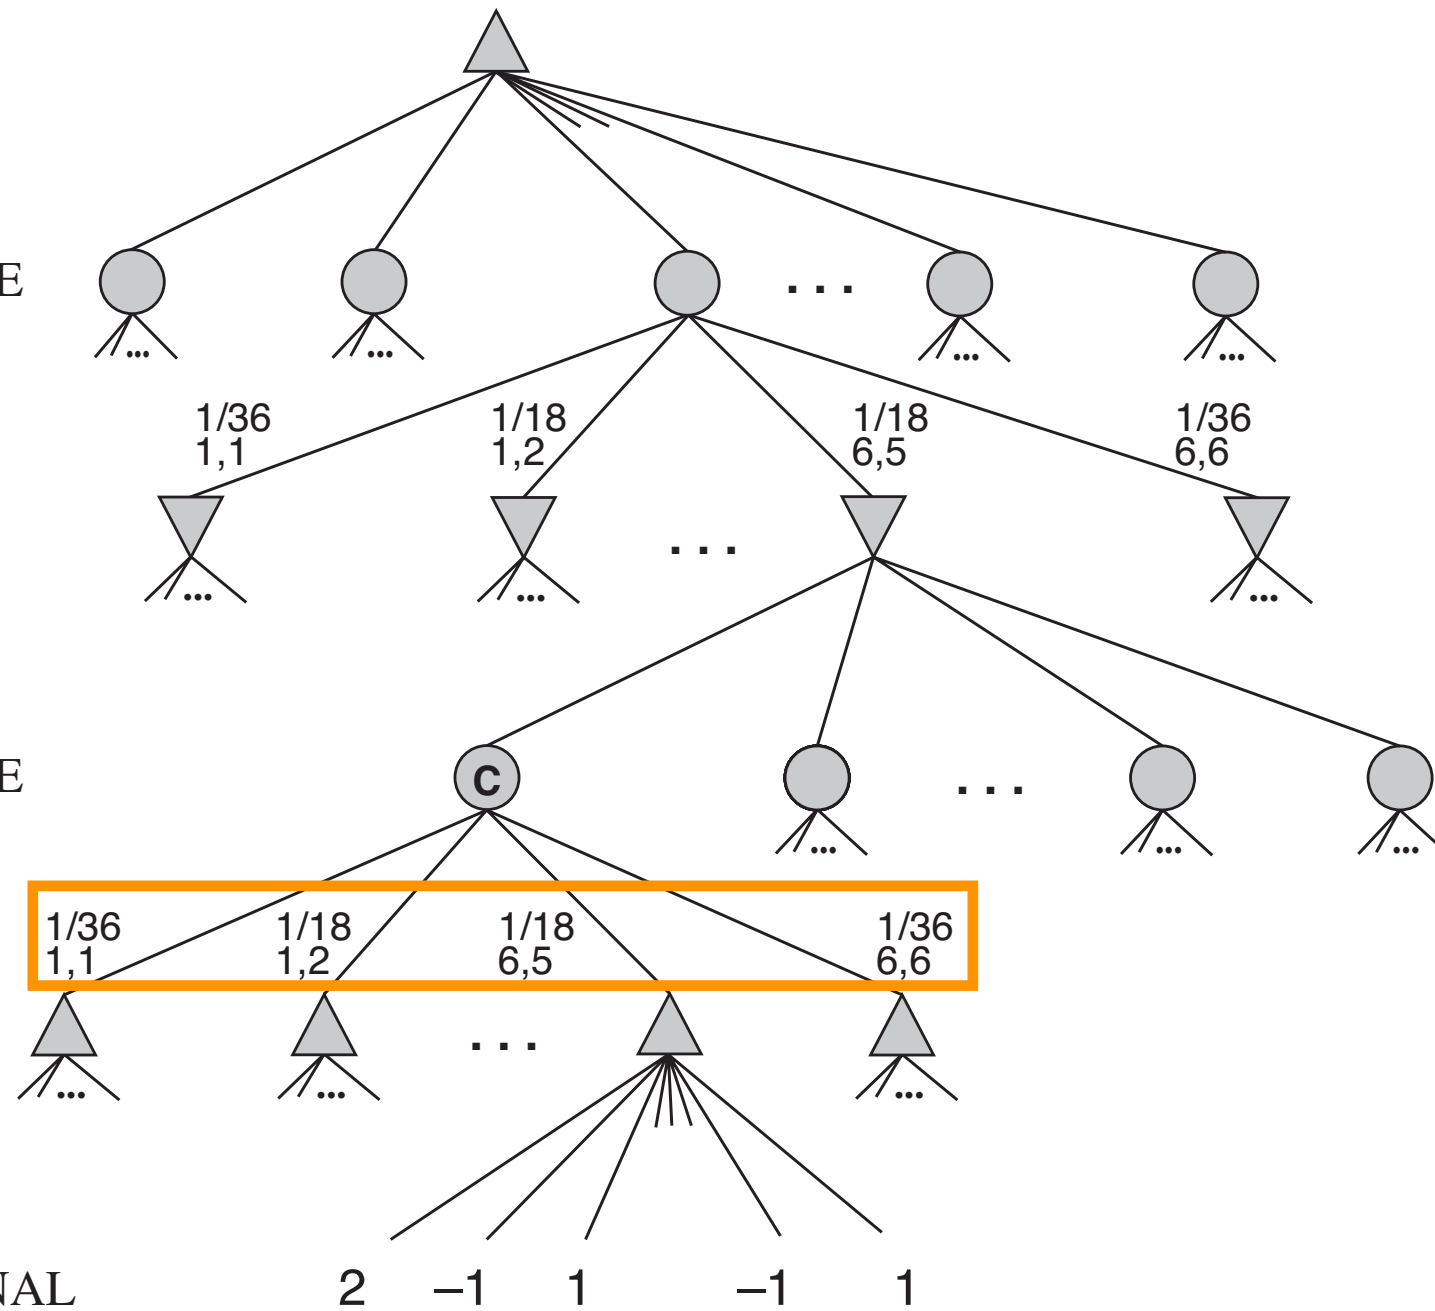
CHANCE

MIN

CHANCE

MAX

TERMINAL





MAX


CHANCE

MIN

CHANCE

MAX

TERMINAL



# Expecti-Minimax

- Weighted average over chance nodes

$$\text{EMINIMAX}(s) =$$

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_a \text{EMINIMAX}(\text{RESULT}(S, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_a \text{EMINIMAX}(\text{RESULT}(S, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \\ \sum_r P(r) \text{EMINIMAX}(\text{RESULT}(S, r)) & \text{if } \text{PLAYER}(s) = \text{CHANCE} \end{cases}$$



MAX

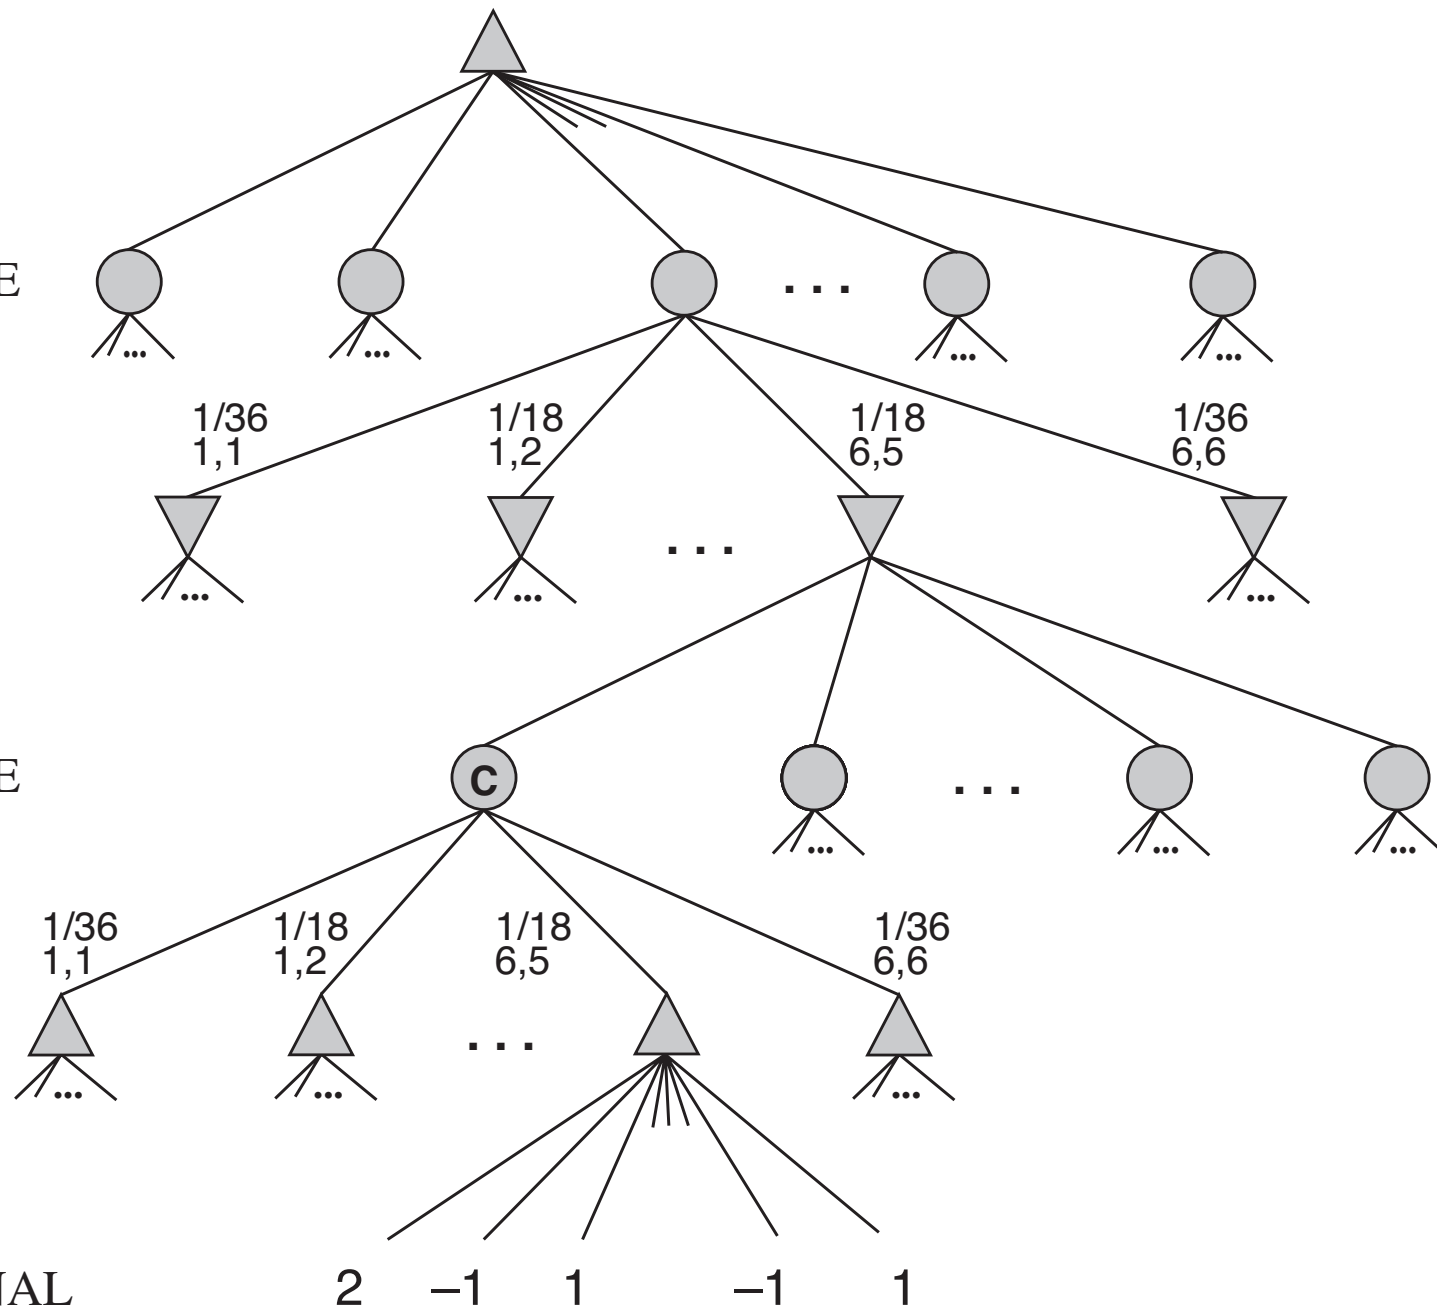
CHANCE

MIN

CHANCE

MAX

TERMINAL



# Expecti-Minimax

- Weighted average over chance nodes

$$O(b^m n^m)$$



# Stochastic Games

- Expectation to handle uncertainty and randomness
- For example in poker: "Pot Odds"

# Partial Observability

- Some of the state of the world is hidden (unobservable)
- There is some uncertainty about the state of the world



# Partially-Observable Games

- Some of the state of the game is hidden from the player(s)
- Interesting because:
  - Valuable real-world games (e.g., cards)
  - Partial observability arises all the time in real-world problems

# Partially-Observable Games

- Deterministic partial observability
  - Opponent has hidden state
  - Battleship, Stratego, Kriegspiel



# Partially-Observable Games

- Deterministic partial observability
  - Opponent has hidden state
  - Battleship, Stratego, Kriegspiel
- Stochastic partial observability
  - Missing/hidden information is random
  - Card games: bridge, hearts, poker (most)



# Stochastic Partially Observable Games





Hand	Frequency	Approx. Probability	Approx. Cumulative	Approx. Odds	Mathematical expression of absolute frequency
Royal flush 	4	0.000154%	0.000154%	649,739 : 1	$\binom{4}{1}$
Straight flush (excluding royal flush) 	36	0.00139%	0.00154%	72,192.33 : 1	$\binom{10}{1}\binom{4}{1} - \binom{4}{1}$
Four of a kind 	624	0.0240%	0.0256%	4,164 : 1	$\binom{13}{1}\binom{12}{1}\binom{4}{1}$
Full house 	3,744	0.144%	0.170%	693.2 : 1	$\binom{13}{1}\binom{4}{3}\binom{12}{1}\binom{4}{2}$
Flush (excluding royal flush and straight flush) 	5,108	0.197%	0.367%	507.8 : 1	$\binom{13}{5}\binom{4}{1} - \binom{10}{1}\binom{4}{1}$
Straight (excluding royal flush and straight flush) 	10,200	0.392%	0.76%	253.8 : 1	$\binom{10}{1}\binom{4}{1}^5 - \binom{10}{1}\binom{4}{1}$
Three of a kind 	54,912	2.11%	2.67%	46.3 : 1	$\binom{13}{1}\binom{4}{3}\binom{12}{2}\binom{4}{1}^2$
Two pair 	123,552	4.75%	7.62%	20.03 : 1	$\binom{13}{2}\binom{4}{2}^2\binom{11}{1}\binom{4}{1}$
One pair 	1,098,240	42.3%	49.9%	1.38 : 1	$\binom{13}{1}\binom{4}{2}\binom{12}{3}\binom{4}{1}^3$
No pair / High card 	1,302,540	50.1%	100%	.005 : 1	$\left[\binom{13}{5} - 10\right] \left[\binom{4}{1}^5 - 4\right]$
Total	2,598,960	100%	100%	1 : 1	$\binom{52}{5}$

# Weighted Minimax

- For each possible deal  $s$ :
  - Assume  $s$  is the actual situation
  - Compute Minimax or H-Minimax value of  $s$
  - Weight value by probability of  $s$
- Take move that yields highest expected value over all the possible deals



# Weighted Minimax

$$\operatorname{argmax}_a \sum_s P(s) \operatorname{MINIMAX}(\operatorname{RESULT}(s, a))$$

# Weighted Minimax

$$\operatorname{argmax}_a \sum_s P(s) \operatorname{MINIMAX}(\operatorname{RESULT}(s, a))$$

$$\binom{26}{13} = 10,400,600$$

$$\binom{47}{25} = 1.48338977 \times 10^{13}$$



# Monte Carlo Methods

- Use a “representative” sample to approximate a events from an underlying probability distribution

# Monte Carlo Minimax

$$\operatorname{argmax}_a \frac{1}{N} \sum_{i=1}^N \text{MINIMAX}(\text{RESULT}(s_i, a))$$



# Summary

- Non-deterministic games
  - Expecti-MINIMAX: Compute expected MINIMAX value over chance nodes
- Partially observable games
  - Weighted MINIMAX: Compute expected value over possible hidden states
- Naive approaches impractical

For Next Time:  
Chapter 4.0–4.1