

CSC242: Introduction to Artificial Intelligence

Lecture 1.6

Please put away all electronic devices

Announcements

- Unit 1 Exam: Next class
- Project 1 due that day 1159PM

Systematic Search

- Enumerates paths from initial state
- Records what alternatives have been explored at each point in the path

Good: Systematic \rightarrow Exhaustive

Bad: Exponential time and/or space

Local Search

- Evaluates and modifies a small number of current states
- Does not record history of search (paths, explored set, etc.)


Good: Very little (constant) memory

Bad: May not explore all alternatives


=> Incomplete

Hill-climbing a.k.a. Greedy Local Search

- Move through state space in the direction of increasing value (“uphill”)



State-space landscape



Random Restart Strategy

```
State randomRestart(Problem p) {  
    while (true) {  
        p.setInitialState(new random State);  
        State solution = hillClimb(p);  
        if (p.isGoal(solution)) {  
            return solution;  
        }  
    }  
}
```

Does it work? Yes (but)

How well does it work?

Prob of success = p Expected # of tries = $1/p$
 ≈ 7
 $= 0.14$



```
State randomRestart(Problem p) {  
    while (true) {  
        p.setInitialState(new random State);  
        State solution = hillClimb(p);  
        if (p.isGoal(solution)) {  
            return solution;  
        }  
    }  
}
```

Randomness

Stochastic hill climbing

Randomness in Search

Pure random walk



Complete,
but horribly slow

Greedy local search




Incomplete,
but fast

Simulated Annealing

- Greedy local search (Hill-descending)
- Select states with lower cost
- OR with some probability even if higher cost
- “High temperature”: higher probability
- “Low temperature”: lower probability
- “Cool” according to schedule

Cost



State space

Simulated Annealing

Complete? No.

Optimal? No.

“But if the schedule lowers T slowly enough, simulated annealing will find a global minimum with probability approaching one.”

Local Search (cont.)

Local Search

- Evaluates and modifies a small number of current states
- Does not record history of search

Good: Very little (constant) memory




Bad: May not explore all alternatives

=> Incomplete




Local Search++

- Keep track of k states rather than just one

Parallel Local Search






Parallel Local Search



Sequential: $1/k$

Parallel Local Search




Sequential: $1/k$

Parallel: No extra cost

Local Search++

- Keep track of k states rather than just one
- At each step, generate all successors of all k states ($k \times b$ of them)
- Keep the most promising k of them

Local Search++



Local Beam Search

- Keep track of k states rather than just one
- At each step, generate all successors of all k states ($k \times b$ of them)
- Keep the most promising k of them
- k = "width of the beam"




Local Beam Search




Local Beam Search


- Keep track of k states rather than just one
- At each step, generate all successors of all k states ($k \times b$ of them)
- Keep the most promising k of them
- k = “width of the beam”

Parallel Local Search



Local Beam Search






1986

"And this report just in. ... Apparently, the grass is greener on the other side."


Local Beam Search

- Keep track of k states rather than just one
- At each step, generate all successors of all k states ($k \times b$ of them)
- Keep the most promising k of them
- k = “width of the beam”

Asexual Reproduction




Mitosis




Sexual Reproduction



Meiosis



Recombination



The Multi-Million Copy Bestseller

RICHARD DAWKINS



THE SELFISH GENE

40TH ANNIVERSARY EDITION



OXFORD LANDMARK SCIENCE

Genetic Algorithms


- Start with k random states
- Select pairs of states and have them “mate” to produce “offspring”
- Most fit (highest-scoring) individuals (states) reproduce more often

Genetic Algorithms

- States encoded as “chromosomes”
(linear sequences, a.k.a. strings)

Genetic Algorithms

- States encoded as “chromosomes” (linear sequences, a.k.a. strings)
- During mating:
 - Crossover: swap chunks of state
 - Mutation: randomly change bits of state



Crossover Mutation

Genetic Algorithms


- States encoded as “chromosomes” (linear sequences, a.k.a. strings)
- During mating:
 - Crossover: swap chunks of state
 - Mutation: randomly change bits of state

GF on GAs

- A version of stochastic local beam search with a special way to generate successor states (motivated by a naive biology analogy)
- “Much work remains to be done to identify the conditions under which genetic algorithms perform well.”



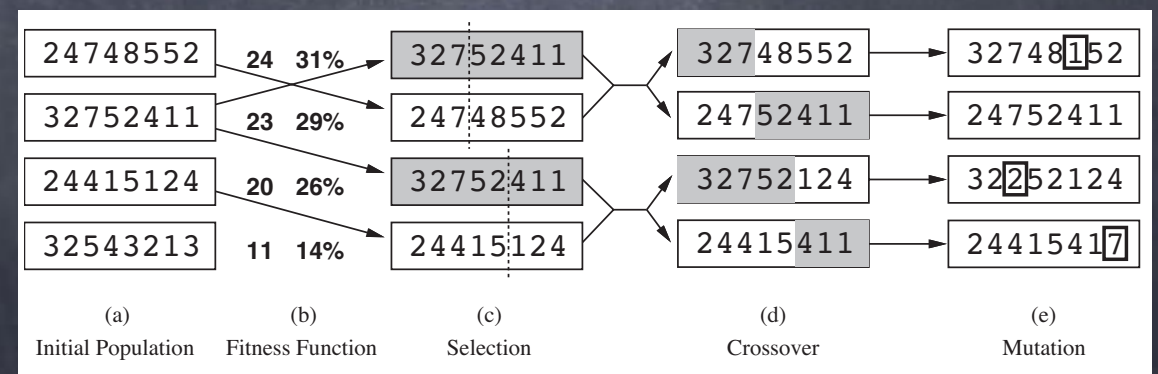
Hill-climbing (Greedy Local Search)



Simulated Annealing



Local Beam Search



Genetic Algorithms


Local Search

- Evaluates and modifies a small number of current states
- Does not record history of search

Good: Very little (constant) memory

Bad: May not explore all alternatives

=> Incomplete



Genetic Algorithms


Local Beam Search

Local Search in Continuous Spaces

- FYI: Section 4.2

Searching with Nondeterministic Actions


Vacuum World




Actions: *Left, Right, Suck*

Vacuum World


1




2




3




4




5




6









7



8




Vacuum World



Fully Observable

Deterministic, Observable Vacuum World



Fully Observable

Nondeterministic Vacuum World

Nondeterministic Vacuum World

- “Erratic vacuum cleaner”: bad *Suck*
 - *Left & Right* work as expected
 - *Suck* on dirty square may clean other square also
 - *Suck* on clean square may make it dirty







Nondeterministic Vacuum World

- “Erratic vacuum cleaner”
- Nondeterministic transition model:
 $\text{RESULT}(s, a)$ returns a set of possible
outcome states

Nondeterministic Vacuum World


- “Erratic vacuum cleaner”
- Nondeterministic transition model:
 $\text{RESULT}(s, a)$ returns a set of possible
outcome states
- Solution is a set of nested if-then
statements (conditional plan)

Vacuum World




Vacuum World


1




2




3




4




5



6









7



8



Vacuum World

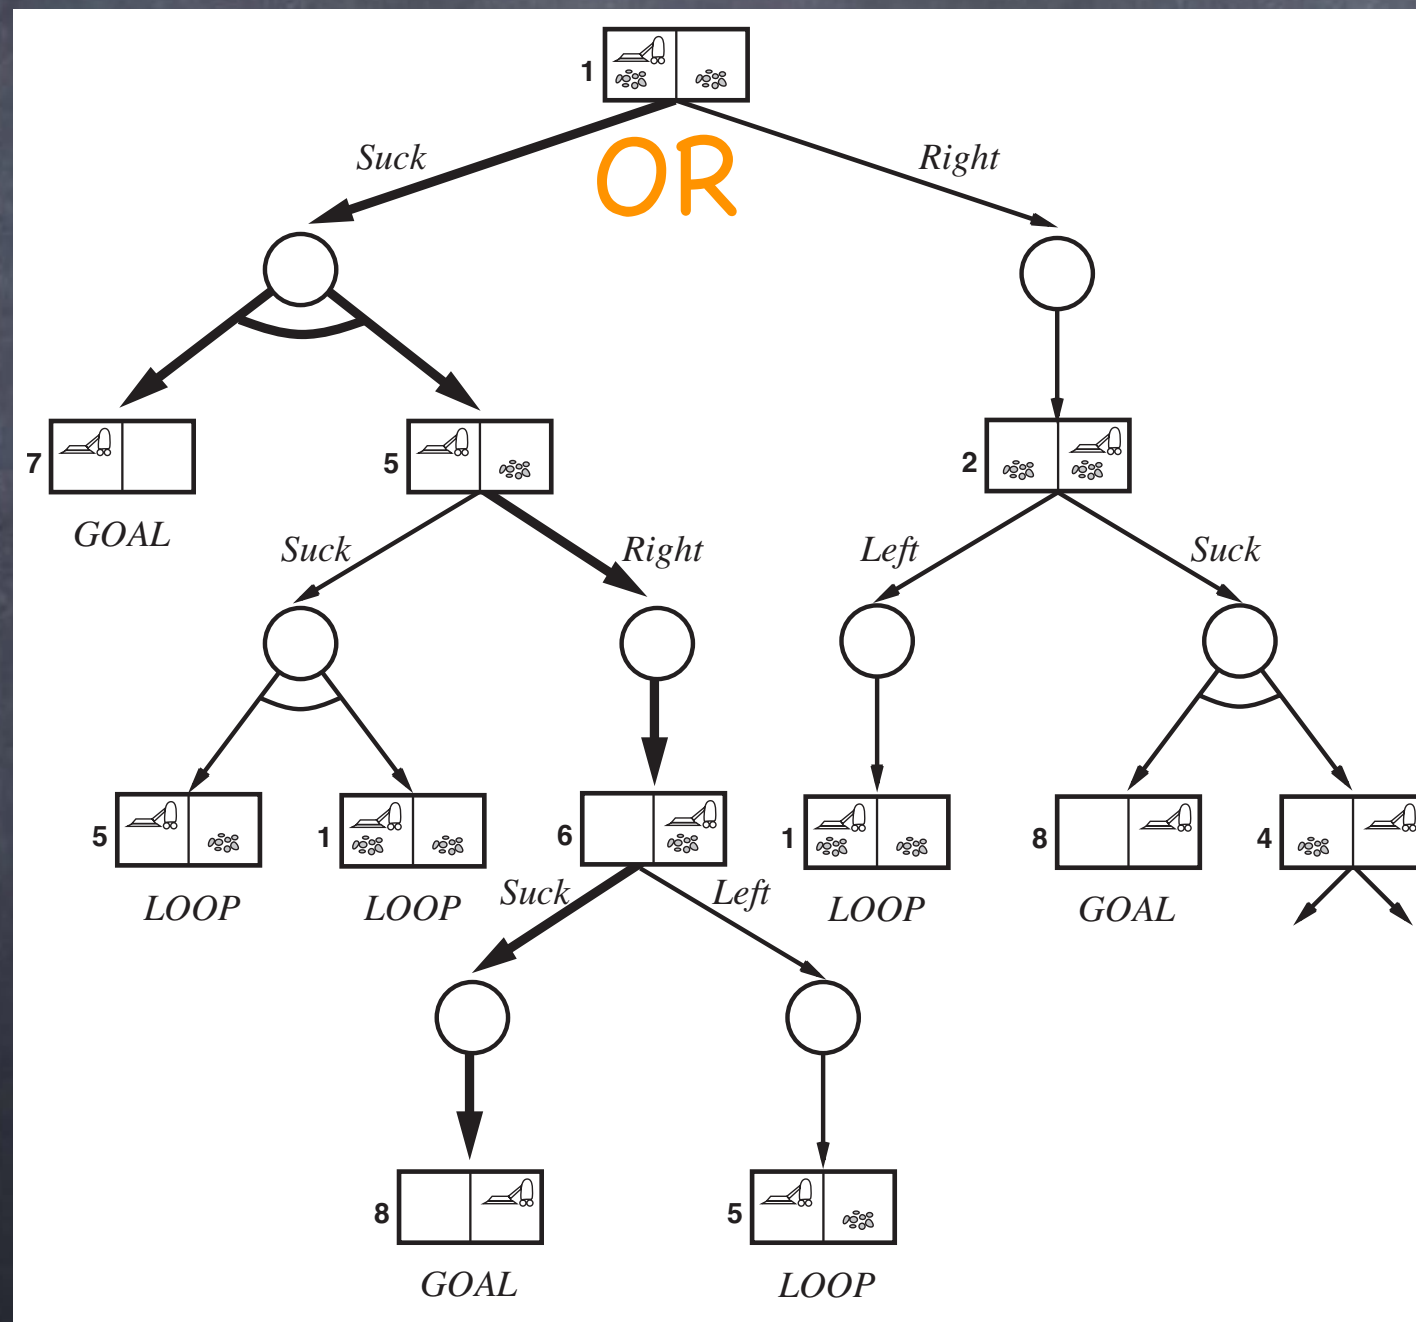


Nondeterministic Vacuum World

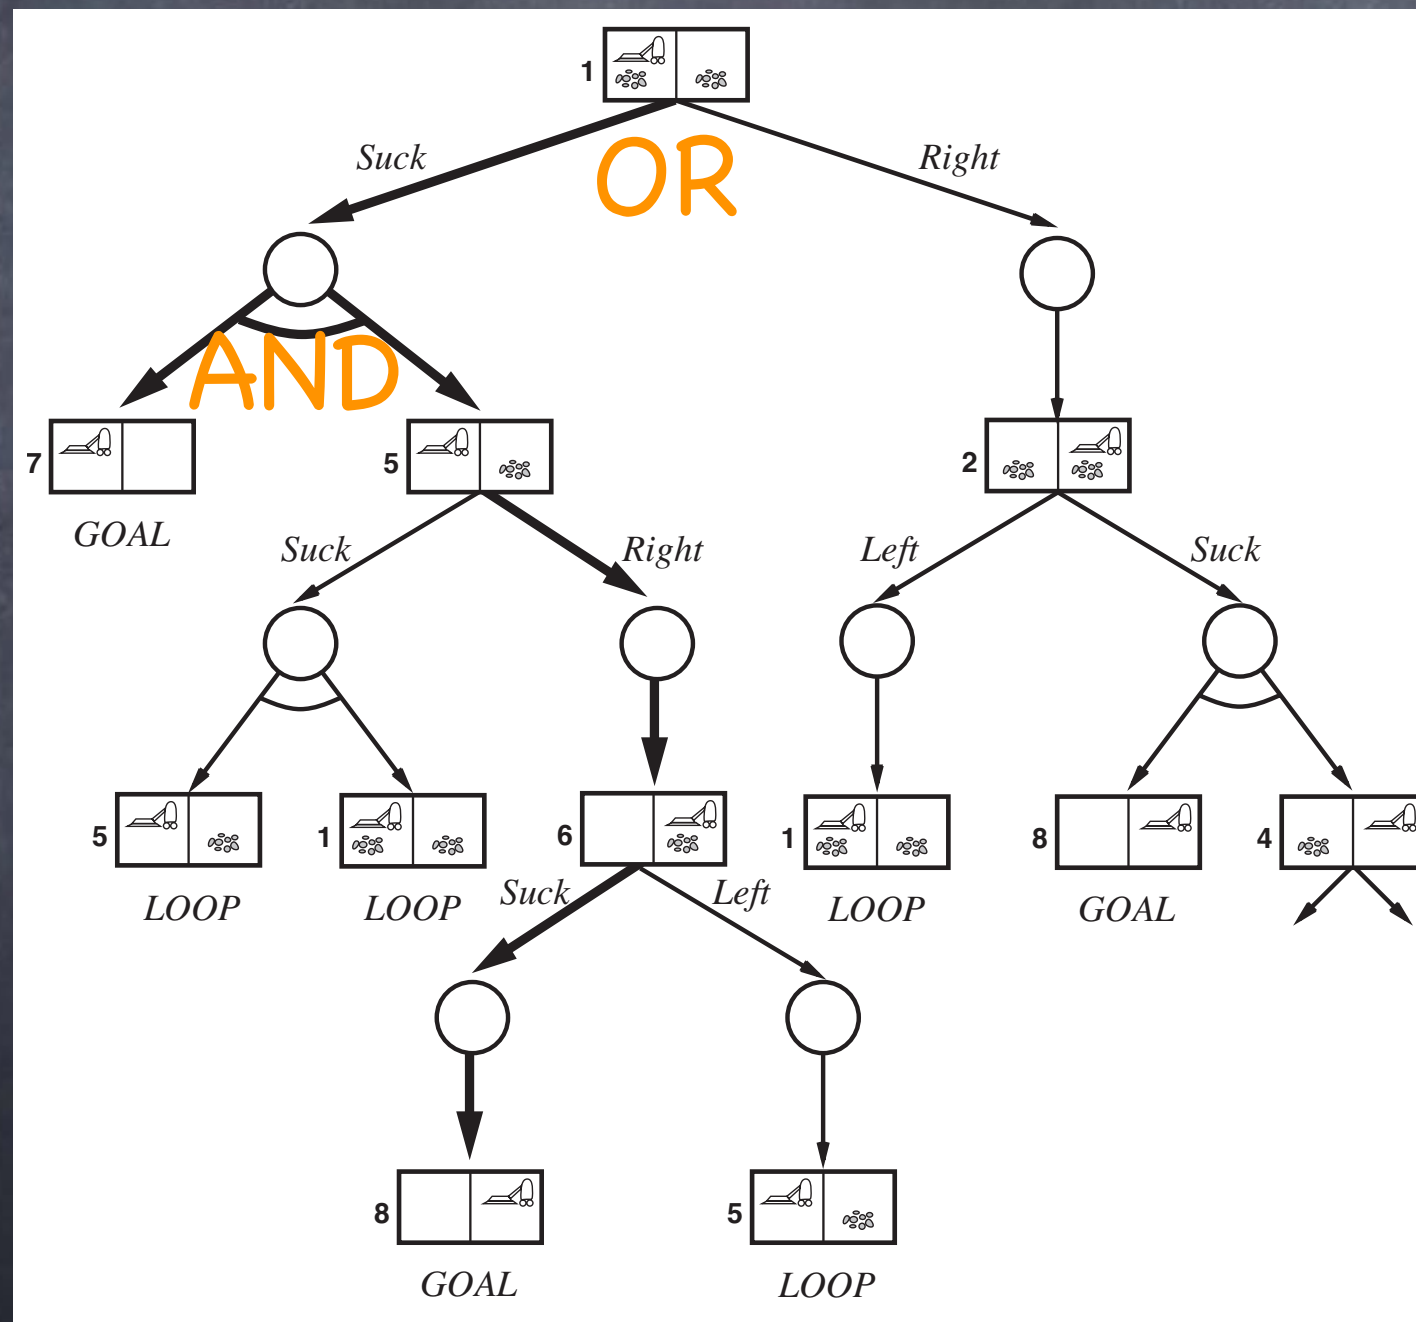
- “Erratic vacuum cleaner”
- Nondeterministic transition model:
 $\text{RESULT}(s, a)$ returns a set of possible outcome states
- Solution is a set of nested if-then statements (conditional plan):

$[Suck, \text{if } State = 5 \text{ then } [Right, Suck] \text{ else } []]$

Nondeterministic Vacuum World

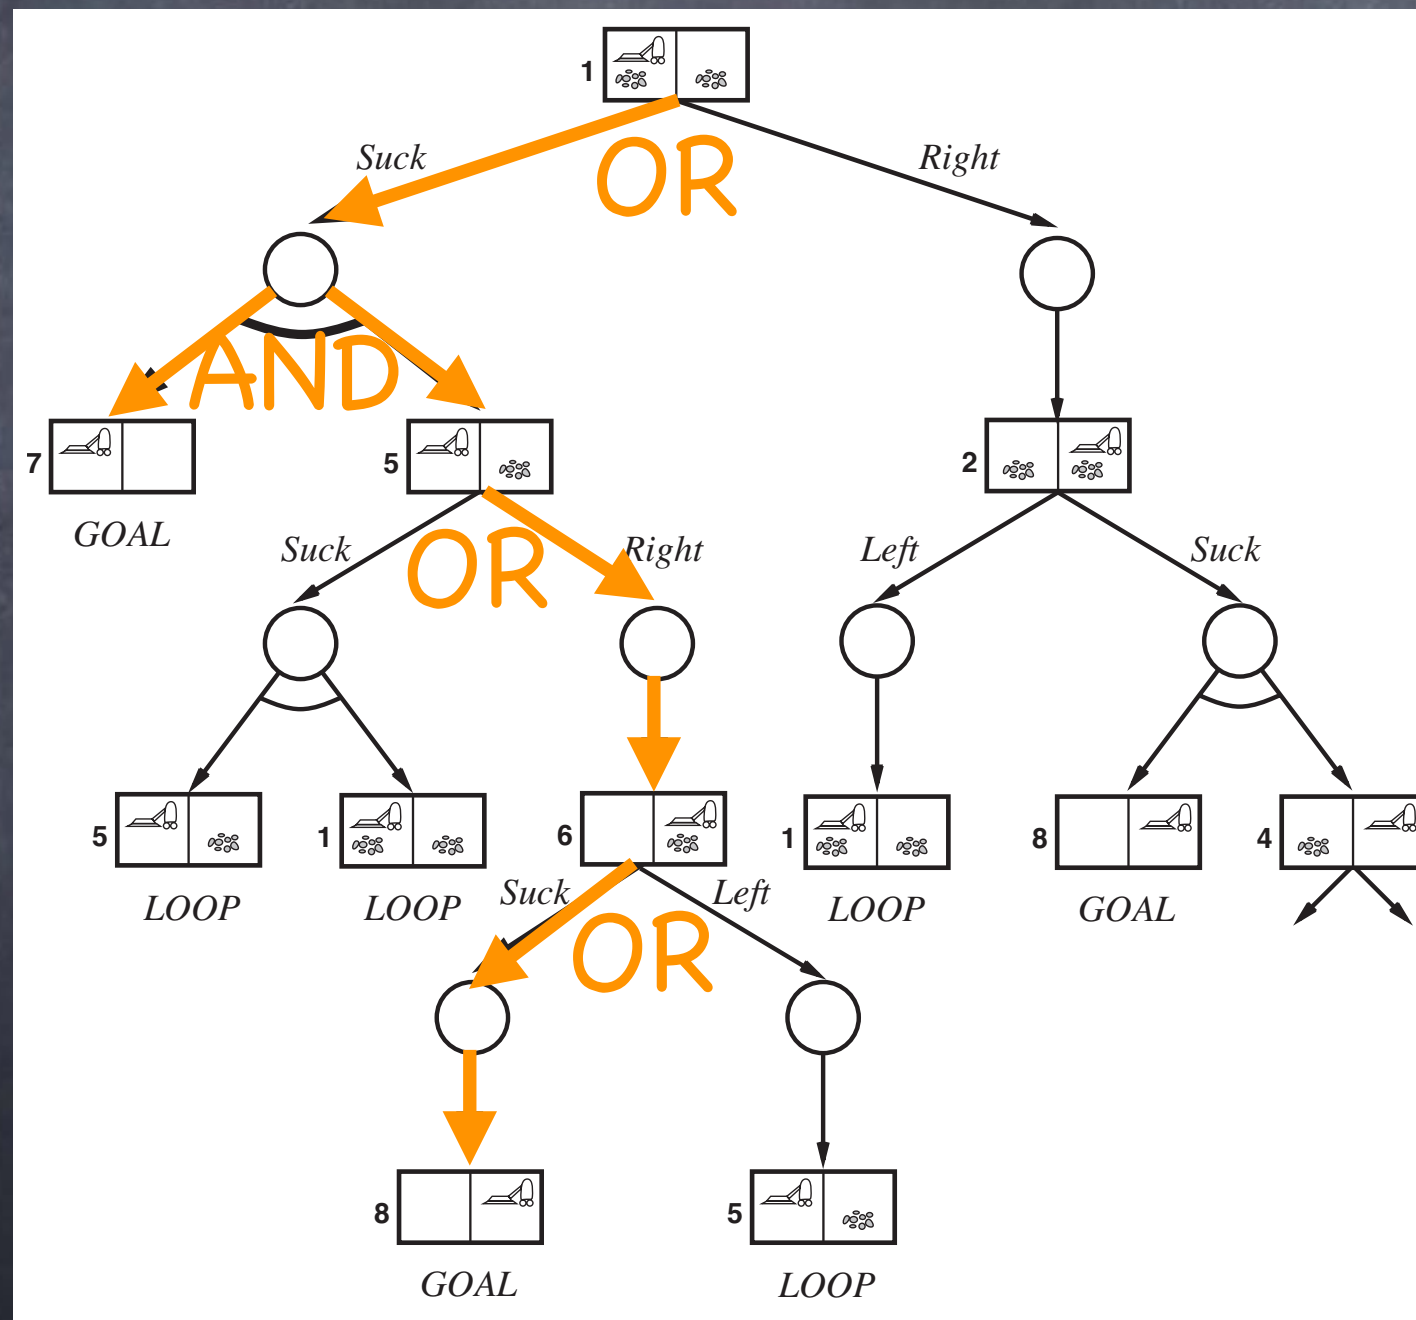


Nondeterministic Vacuum World



The diagram illustrates a search tree for a vacuum cleaner problem. The root node is a state with a vacuum on the left and dirt on both sides, labeled '1'. It branches into 'Suck' and 'Right' actions. The 'Suck' branch leads to a node with a vacuum on the left and dirt on the right, labeled '5', which then branches into 'Suck' (leading to a goal state '7') and 'Right' (leading to a node with a vacuum on the right and dirt on both sides, labeled '6'). The 'Right' branch from the root leads to a node with a vacuum on the right and dirt on both sides, labeled '2', which branches into 'Left' (leading to a node with a vacuum on the left and dirt on both sides, labeled '1') and 'Suck' (leading to a node with a vacuum on the right and dirt on the left, labeled '4'). The 'Left' branch from node '6' leads to a node with a vacuum on the left and dirt on the right, labeled '5', which branches into 'Suck' (leading to a goal state '8') and 'Right' (leading to a node with a vacuum on the right and dirt on both sides, labeled '5'). The 'Right' branch from node '4' leads to a node with a vacuum on the right and dirt on the left, labeled '4', which branches into 'Suck' (leading to a goal state '8') and 'Right' (leading to a node with a vacuum on the right and dirt on the left, labeled '4'). The diagram includes labels 'GOAL' and 'LOOP' for various states and uses orange text 'AND' and 'OR' to indicate logical relationships between branches.

Nondeterministic Vacuum World









Searching with Nondeterministic Actions

- Depth-first algorithm: AIMA Fig. 4.11
- Also breadth-first, best-first, incl. A^*








Nondeterministic Vacuum World

- “Slippery vacuum world”:
 - bad *Left/Right*
- Transition model:
 - $\text{RESULT}(s, a)$ may include s

Vacuum World



Vacuum World



Nondeterministic Vacuum World

- “Erratic vacuum cleaner”: bad *Left* and/or *Right*
- Transition model: $\text{RESULT}(s, a)$ may include s
- Solution requires loops


Searching with Nondeterministic Actions

- Modified version of AND-OR graph search (AIMA Fig. 4.11)
- Will find cyclic solutions
 - “Provided that each outcome of a nondeterministic action eventually occurs.”


Searching with Partial Observations

Sensorless Vacuum World


1




2




3




4




5




6




7



8




Sensorless Vacuum World



Sensorless Vacuum World


Right




Sensorless Vacuum World

Right


1




2




3




4




5




6



7




8




Sensorless Vacuum World

*Right
Suck*


1




2




3




4




5




6



7




8




Sensorless Vacuum World

*Right
Suck*


1




2




3




4




5




6



7









8




Sensorless Vacuum World

*Right
Suck
Left*




Sensorless Vacuum World

*Right
Suck
Left*




Sensorless Vacuum World

*Right
Suck
Left
Suck*




Sensorless Vacuum World

*Right
Suck
Left
Suck*



Sensorless Vacuum World

*Right
Suck
Left
Suck*



Belief State

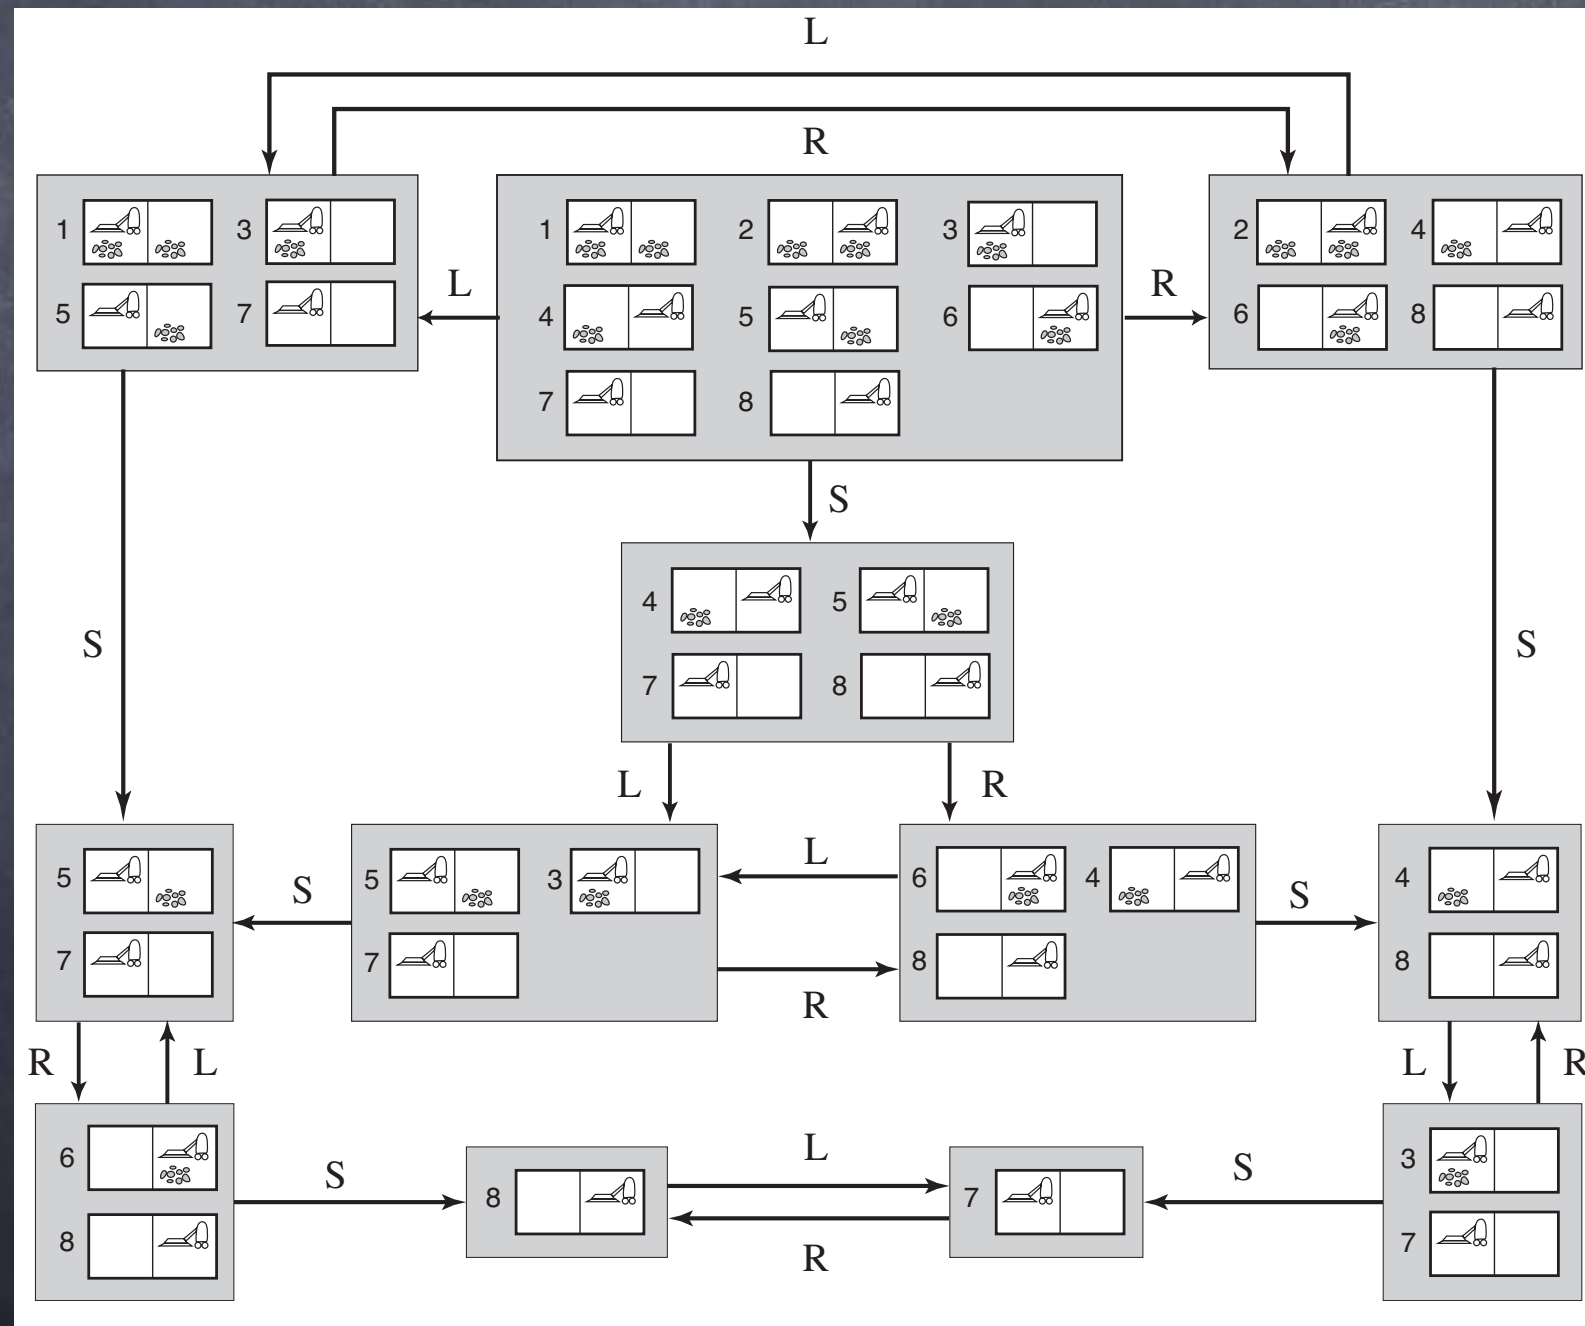
- An agent's current beliefs about what physical (real-world) state it might be in
- Belief state is fully observable
 - The agent knows its own beliefs (for now)

Belief State-Space Search

- Belief states: Every possible set of physical states (N physical, 2^N belief)
- Initial state: Set of all states in problem
- Actions: Union or intersection of applicable actions in each state in the belief state
- Transition model: Union (liberal) or intersection (conservative) of result states (prediction)
- Goal test: Every state in belief state is a goal state
- Path cost: "Tricky"

Vacuum World

Belief State Space



Searching with Partial Observations

- Search through space of belief states
- Sensorless problem-solving “completely impractical”
- Solutions:
 - Use a better representation of state
 - Incorporate observations that update the current belief state

Summary

- Local Search
 - Hill-climbing, Random Restarts, Simulated Annealing, Local Beam Search, Genetic Algorithms
- Search in nondeterministic domains
 - AND-OR trees, contingency plans
- Search in partially-observable domains
 - Belief states, belief-state-space search

Unit 1: Search

- Problem solving and state-space search
- Search strategies
- Adversarial search
- Local search
- Search in nondeterministic and partially observable domains

For next time:

Unit 1 Exam